

Operating system structure

1- Monolithic structure -Original UNIX

Two parts:

System programs

The kernel

Benefits and limitations:

- +Performance → Efficient from communication side of view
- +little overhead
- +speed
- Debugging → Huge line of code (thousands)
- Difficult to implement and extend
- Difficult to change → tightly coupled

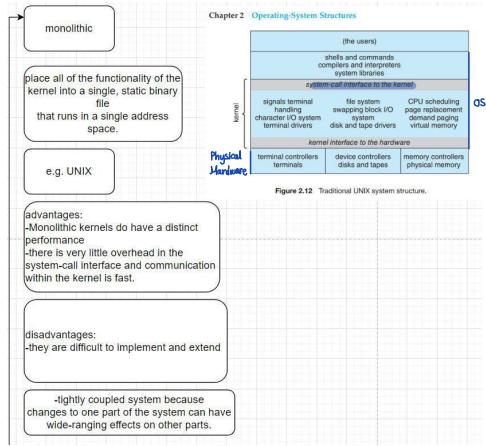


Figure 2.12 Traditional UNIX system structure.

3-Microkernel

Moves from kernel into user space Only small piece is still in kernel

Message passing

Example: Mach and Darwin

Benefits and limitations:

- +More reliable, secure
- +Easier to extend
- +Easier to port OS to new architecture
- Performance overhead of user and kernel communication
- System function overhead

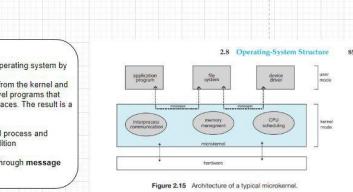
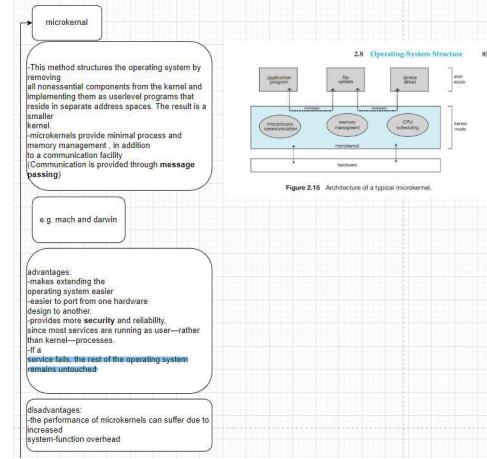


Figure 2.15 Architecture of a typical microkernel.

4-Modules The best

bootable kernel Modules

Object-oriented approach

Each core component is separate and talk through interfaces

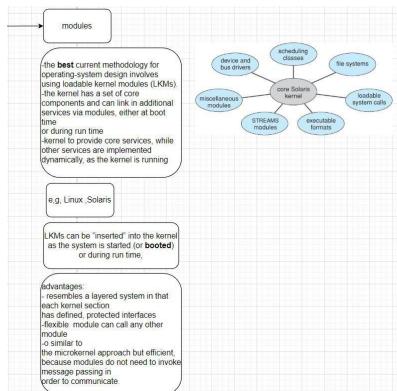
Similar to layers but more flexible

No message passing

Example: Linux, Solaris

Benefits and limitations:

- +Efficient from communication side of view
- +Protection
- +Flexible



5-Hybrid systems

linux and solaris. Monolithic+modular

Windows: monolithic+microkernel

Apple Mac OS X hybrid layered Aqua UI, Cocoa programming environment

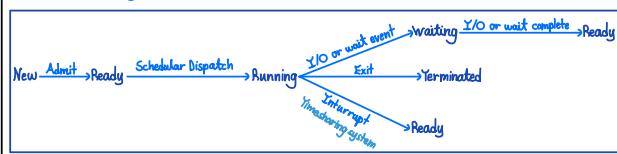
Process

Process address space, memory image parts:



Process state:

- 1-New
- 2-Ready
Ready for execution but currently not under execution
- 3-Running
- 4-Waiting
- 5-Terminated
Completed normally or crashed



Process Control Block(PCB), Task Control Block:

- 1-Process state
- 2-Process number
- 3-Program Counter Register-PC
Points to the next instruction to be executed
- 4-CPU Registers
Instruction Register-IR
Includes the one instruction under execution
IR size: one word
- 5-CPU scheduling information
- 6-Memory management information
- 7-Accounting information
- 8-X/O status information

Context switch:

Saving the old process state(Screenshot), load new process saved state
Pure overhead

Process scheduling:

Multitasking process should be scheduled → Multiprocess → one CPU

Scheduling queue:

- 1-Job queue
longer than ready queue
- In Hard disk, since Main memory is not enough

2-Ready queue

- In Main memory
- 3-Devices queues

Schedulers:

- 1-Short-term schedulers
CPU scheduler
Invoked frequently → ms

2-long-term schedulers

- Job scheduler
Invoked infrequently → s/m
Control degree of multiprogramming
Good process mix

3-Medium-term schedulers

- If degree of multiprogramming needs to decrease
Swapping Memory ←→ Disk

Process types:

- 1-I/O-bound process
Many-Short CPU bursts
- 2-CPU-bound process
Few-Very long CPU bursts

System mechanism for process:

1-Process creation

Resource options:

- 1-Parent and child share all resources
- 2-Child share subset of parent resources
- 3-Parent and child share no resources

Execution options:

- 1-Parent and child execute concurrently
- 2-Parent wait until child terminates

Address space:

- 1-Child duplicate of parent

System will create an address space for child in a different space than parent's space
Child address space is a copy paste from parent's address space

fork() return value:

- If successful → will return new process child ID > 0
If failure error → will return < 0
I am child → will return 0

2-Child has program loaded into

System will create an address space for child in a different space than parent's space
Child address space is a copy paste from parent's address space
New program is loaded into child

Methods:

Linux-Paxos: C-fork() Creation
exec() exec() execvp() handing
Windows: CreateProcess() Creation+handing

2-Process termination

exit() - Process terminate itself

Return status data from child to parent and pid using wait()

Process resources deallocated by OS

abort() - Parent process terminate child process

Reasons:

- 1-Child exceeded allocated resources
- 2-Task assigned to child no longer required
- 3-The parent is exiting and OS doesn't allow child to continue

Cascading termination

All child and grandchild are terminated

Orphan

If parent terminated without invoking wait

3-Interprocess communication(IPC)

1-Shared memory

Cooperating process should be in the same device → Access to common memory

Frequent messaging, large amount of data

Yester as there is no system calls → One system call

No kernel interference

Type of buffers:

- 1-Unbounded
- 2-Bounded

2-Message passing

Cooperating process in the same device or different devices

System calls for every sending and receiving

Infrequent messaging, small amount of data

Communication link+Method

Methods:

- 1-send(message)
- 2-receive(message)

Communication link implementation:

1-Physical

- 1-Shared memory
- 2-Hardware bus
- 3-Network

2-logical

1-Direct or Indirect

Indirect methods: create and destroy or delete port, send and receive

2-Synchronous or Asynchronous

If both are synchronous → rendezvous

3-Automatic or explicit buffering

- 1-Zero capacity
- 2-Bounded capacity
- 3-Unbounded capacity

Process types:

1-Independent

2-Cooperating

Advantages:

- 1-Information sharing
- 2-Computation speedup
- 3-Modularity
- 4-Convenience

Operating system execute programs:

- 1-Batch system - jobs
- 2-Time-shared system - User program or task

Threads

Kernel are generally multithreaded
Basic unit of CPU utilization

Each Thread has:

- 1-Thread ID
- 2-PC
- 3-Register set
- 4-Stack

Threads share:

- 1-Text section
- 2-Data section
- 3-OS resources

Threads benefits:

- 1-Responsiveness
- 2-Resource sharing
- 3-Economy
- 4-Scaleability

Computer system architecture:

- 1-Single processor systems
 - One chip one core
- 2-Multiprocessors-parallel systems
 - Many chips each has core
- 3-Multicore
 - One chip many core's

Multiprocessors and Multicore challenges:

- 1-Dividing activities
- 2-Balance
- 3-Data splitting
- 4-Data dependency
- 5-Testing and debugging

Parallelism and concurrency:

Concurrency:

Single processor

Parallelism → Simultaneously:

Multiprocessors and Multicore

1-Data parallelism

Dividing data into cores

All cores are performing same operation

Working on same code and function, but different sections

2-Task parallelism

Every core is performing different operation (Unique) on same data

Amdahl's law:

$$\text{speedup} \leq \frac{1}{S + \frac{(1-S)}{N}}$$

Thread creation categories:

- 1-Synchronous
- 2-Aynchronous

Multithreading supported by:

- 1-User-level libraries - User threads
 - Unknown to kernel
 - Faster than kernel threads
- 2-kernel itself
 - Eg: Operating systems

Multithreading models:

- 1-Many-to-One

User thread
Efficient but didn't get full benefits of process
Not popular

Examples:
Solaris green, GNU Portable

- No need for kernel support for multithreading (+)
- Thread creation is fast (+)
- Switching between threads is fast; efficient approach (+)
- Blocking system calls defeat the purpose and have to be handled (-)
- Multiple threads will run on a single processor, not utilizing multi-processor machines. (-)

- 2-One-to-One

kernel thread
Popular, most common

Examples:
Windows, Linux

- Provides more concurrency; when a thread blocks, another can run. Blocking system calls are not problem anymore. (+)
- Multiple processors can be utilized as well. (+).
- Need system calls to create threads and this takes time; (-)
- Thread switching costly; Number of threads per process sometimes restricted due to overhead (-)
- any thread function requires a system call. (-)

- 3-Many-to-Many

Between the two previous
Not so successful
User threads => kernel threads

- 4-Two level

Similar to Many-to-Many, but allow user to bound kernel

Thread main libraries:

Provide API

- 1-Pthread

Either user or kernel level

POSIX

Specification not implementation

Data declared globally

Common:

UNIX, Linux, MAC OS X

Methods:

`pthread_attr_init()` Default initialization
`pthread_create(4 parameters)` Creation, Return 0 if succeed

`pthread_exit()`

`pthread_join()`

`pthread_cancel()`

- 2-Windows

kernel level

Data declared globally

- 3-Java, thread API

CPU Scheduling

Ready Queue

Smallest entity to be scheduled in CPU threads

CPU decisions when:

- 1-Switch Running → Waiting Nonpreemptive
Eg: I/O request, wait)
- 2-Switch Running → Ready Preemptive
Eg: Interrupt, timer interrupt
- 3-Switch Waiting → Ready Preemptive
Eg: I/O completion
- 4-Terminates Nonpreemptive

Nonpreemptive and Preemptive:

Nonpreemptive:

Process makes the CPU

Preemptive:

All OS uses

Consider:

- 1-Shared data
- 2-Preemption while in kernel time
Affect design of OS kernel!
- 3-Interrupts occurring during crucial OS activities

Dispatcher:

- 1-Context switch < 10 microseconds
 - 1-Voluntarily
 - 2-Involuntarily
- 2-Switch to user mode
- 3-Jump to the proper location in user program to restart the program

Scheduling criteria:

- 1-CPU Utilization MAX

Range 40% - 90%

$$\text{CPU Utilization} = \frac{\text{Active time CPU}}{\text{All time CPU}} = 1 - \frac{\text{Total CPU idle time}}{\text{All time CPU}}$$
- 2-Throughput MAX

Depend on CPU and process types:

 - 1-long process low throughput
 - 2-Short process High throughput
- 3-Turnaround time MIN

Completion time - Arrival time
- 4-Waiting time MIN

Turnaround time - Burst time
- 5-Response time MIN

First response - Arrived time

Scheduling Algorithms:

- 1-First Come First Serve (FCFS): Nonpreemptive Single processor, CPU, Core

Special case of priority

Problem: Average waiting time

Trouble some to interactive systems

Convey effect: long process then short process

One CPU-bound, Many I/O-bound process

Lower CPU and device utilization

- 2-Shortest Job First (SJF): Nonpreemptive Single processor, CPU, Core

Special case of priority

Shortest next CPU bursts

Optimal—Gives minimum average waiting time

Estimate length of next CPU burst (Exponential average):

$$T_{n+1} = \alpha T_n + (1-\alpha) T_n$$

$\alpha=0$

• $T_{n+1} = T_n$

• Recent history does not count

$\alpha=1$

• $T_{n+1} = \alpha T_n$

• Only the actual last CPU burst counts Recent history

$\alpha=0.5$

Recent and past history has equal weight

- 3-Shortest Remaining Time First: Preemptive version of SJF Single processor, CPU, Core

- 4-Round Robin (RR): Preemptive Single processor, CPU, Core

Treated as FIFO

Higher turnaround than SJF but better response

n process in ready queue → Each will have $\frac{1}{n}$ CPU time

No process waits more than $(n-1)q$

Time quantum-Time slice (q):

Usually 10-100 ms

q large → FIFO

q small → must be large since context switch will be overhead if q is small.

- 5-Priority scheduling: Preemptive or Nonpreemptive Single processor, CPU, Core

Priority:

- 1-Internally defined

Based on OS

- 2-Externally defined

Process comes with its number

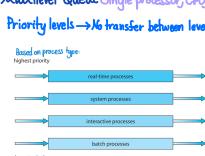
Eg: Based on importance, type and amount

Problem: Starvation

Solution: Aging

- 6-Multilevel Queue Single processor, CPU, Core

Priority levels → No transfer between levels



- 7-Multilevel Feedback Queue Single processor, CPU, Core

Multilevel-feedback-queue scheduler defined by the following parameters:

- number of queues
- scheduling algorithms for each queue
- method used to determine when to upgrade a process
- method used to determine when to demote a process
- method used to determine which queue a process will enter when that process needs service

- 8-Multiple processor scheduling

- 1-Heterogeneous-Asymmetric Master server

- 2-Homogeneous-Symmetric
 - More common since it's distributed
 - load balancing

Synchronization tools

Critical section:
Code segment that access shared data and has to be executed in atomic action

Critical section problem solution(Conditions):

- 1-Mutual exclusion
If process i is executing in critical section, no other process can be executing in critical section
- 2-Progress
Selection of next process to enter critical section cannot be postponed
- 3-Bounded waiting
Bound the number of times that the process have asked to wait

Critical section handling in OS:

- 1-Preemptive
Allow process to be preempted will running in kernel mode
More responsive → Suitable for real-time programming
- 2-Non-Preemptive
Doesn't allow process to be preempted will running in kernel mode
Free from race conditions

Software solutions:

- 1-Global int variable
 - 1-Mutual exclusion yes
 - 2-Progress no
 - 3-Bounded waiting yes
- 2-Two global boolean variables
 - 1-Mutual exclusion no
 - 2-Progress yes
 - 3-Bounded waiting yes
- 3-Two global boolean variables → Turn flag before entering loop
 - 1-Mutual exclusion yes
 - 2-Progress no
 - 3-Bounded waiting yes

Peterson's solution:

Software-based solution to critical section problem
hood, store machine language must be atomic

Not guaranteed to work on modern architecture

Share two variables:

- 1-int turn
- 2-boolean flag[2]

Hardware-based solution to critical section problem:

Uniprocessor to disable interrupt → Not efficient on multiprocessor systems

High-level software tools solution to critical section problem:

- 1-Mutex (Mutual exclusion)
 - Simplist
 - Using locks
 - Variables: boolean available
 - Methods: acquire(), release() atomic
 - Disadvantage: Busy waiting = Spinlock No context switch
- 2-Semaphore
 - More robust tool
 - Provide more sophisticated ways for process to synchronize
 - Variables: int S, list of process
 - Methods: wait() → P(), signal() → V()
 - Semaphore value:
 - 1-Counting
 - Unrestricted domain
 - 2-Binary
 - Between 0 and 1
 - Same as mutex

Semaphore could be with busy waiting or without busy waiting

With busy waiting → waiting queue:

- FIFO To insure bounded waiting
- Data items: int value + pointer to next in list
- Operations:
 - block() or sleep() Place into waiting queue
 - wakeup() Waiting state → ready state

liveness:

liveness failure:

Indefinite loop

Deadlock

Classical problems of synchronization:

- 1-Bounded buffer
 - Shared data:
 - Semaphore mutex → 1
 - Semaphore full → n
 - Semaphore empty → n
- 2-Readers and writers
 - Reader → only read
 - Writer → read and write
 - Shared data:
 - Data set
 - Semaphore rw_mutex → 1
 - Semaphore mutex → 1
 - int read_count → 0

First variation: No reader wait unless writer is writing

Second variation: Once writer is ready → write ASAP

Solution:

Readers-writers locks

- 3-Dining Philosophers

Shared data:

Data set

Semaphore chopstick[5] → 1

Possible solutions:

- instead of sleeping for a fixed time, the philosopher would sleep for some random amount of time (i.e sleep(random_time)).
- Allow at most four philosophers to be sitting simultaneously at the table.
- Allow a philosopher to pick up her chopsticks only if both chopsticks are available (to do this, she must pick them up in a critical section).
- Use an asymmetric solution—that is, an odd-numbered philosopher picks up first her left chopstick and then her right chopstick, whereas an even-numbered philosopher picks up her right chopstick and then her left chopstick.

Main memory

Memory unit only see:
 1-Address, read request
 2-Address, data and write request

Register → One CPU clock or less
 Main memory → Many cycles Causing a stall.

Protection provided by:
 1-Base (base register=relocation register)
 Hold smallest legal physical address
 2-limit registers

Address representation different stages:
 1-Symbolic
 Source code
 Variable content
 2-Relocatable
 Compiled code
 X bytes from the beginning of module → logical.
 3-Absolute
 linker or loader
 Address → Physical.

Address binding:
 1-Compile time
 2-link time
 3-Execution time

Dynamic loading:
 No special support from OS
 Can help by providing libraries to implement dynamic loading
 Implemented through program design

Linking:
 1-Static linking
 Combined by loader
 2-Dynamic linking

Memory management schemes:

Contiguous allocation:
 Old method
 1-Fixed partitioning
 Variable sizes but will never change
 Internal fragmentation+External fragmentation

2-Variable partitioning

OS should maintain:
 1-Allocated partitions
 2-Free partitions→hole

External fragmentation

Non-Contiguous allocation:

1-Paging

Frames:
 Fixed size → power of 2
 Physical memory

Pages:

Fixed size
 Logical memory
 Avoid external fragmentation
 Still have internal fragmentation

2-Segmentation

Segments:
 Variable size

Fragmentation:

1-Internal fragmentation
 2-External fragmentation
 Reduced by compaction

Dynamic storage allocation:

1-First-Fit
 N blocks allocated → 0.5 N blocks lost to fragmentation → 1/3 unusable → 20 percent rule
 2-Best-Fit
 3-Worst-Fit

Virtual memory

Benefits:
 More concurrent programs
 less I/O needed to swap/load portions of programs
 Execute programs larger than RAM size
 Easy sharing of address spaces by several process
 Allow more efficient process creation

Demand pages:
 less I/O needed
 less memory needed
 Faster response
 More users

Page fault activities:
 1-Service the interrupt
 2-Read the page
 3-Restart the process

Page fault rate:
 $\text{Page fault rate} = \frac{\text{Page faults}}{\text{Total memory references}}$

0 → no page fault
 1 → Every reference is page fault

Effective access time:

$$EAT = (1-p) \times \text{memory access} + p \times (\text{page fault overhead} + \text{swap page out} + \text{swap page in})$$

What if no free frames:
 1-Borrow options
 2-Prevent over allocation
 3-Page replacement

Page replacement algorithm:

- 1-FIFO algorithm 3
- 2-Optimal algorithm 1
- 3-Recently used 2
- 4-Highest frequently used Counting algorithm
- 5-Most frequently used Counting algorithm

Frame allocation:

1-Equal allocation
 2-Proportional allocation

$$\alpha_i = \frac{s_i}{\sum s_i} \times \text{total number of frames}$$

Frame allocation:

- 1-Global allocation
- 2-Local allocation